

# DialEgg

## Dialect-Agnostic MLIR Optimizer using Egglog

*EGRAPHS Community Meeting  
August 21, 2025*

**Abd-El-Aziz Zayed and Christophe Dubach**



McGill

# Equality Saturation in the CASL Lab



**PI Christophe Dubach**



**Jonathan Van der Cruysse**

Latent idiom recognition with a minimalist IR using EqSat



**Tzung-Han Juang**

EqSat for hardware resource sharing



**Adam Musa**

Driving EqSat with RL

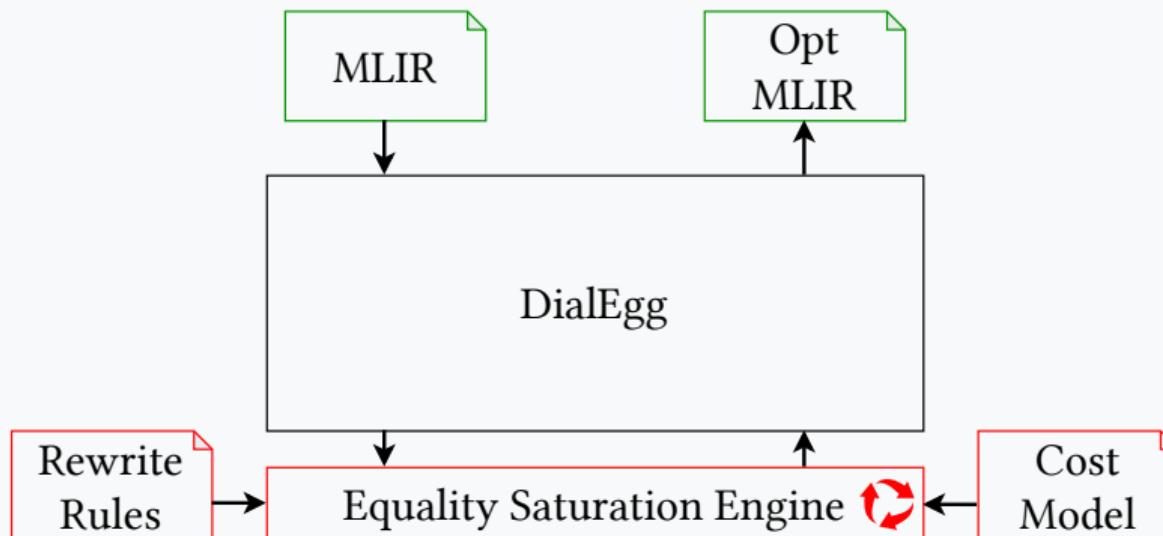


**Aziz Zayed (me)**

EqSat and egraphs in modern compilers like MLIR

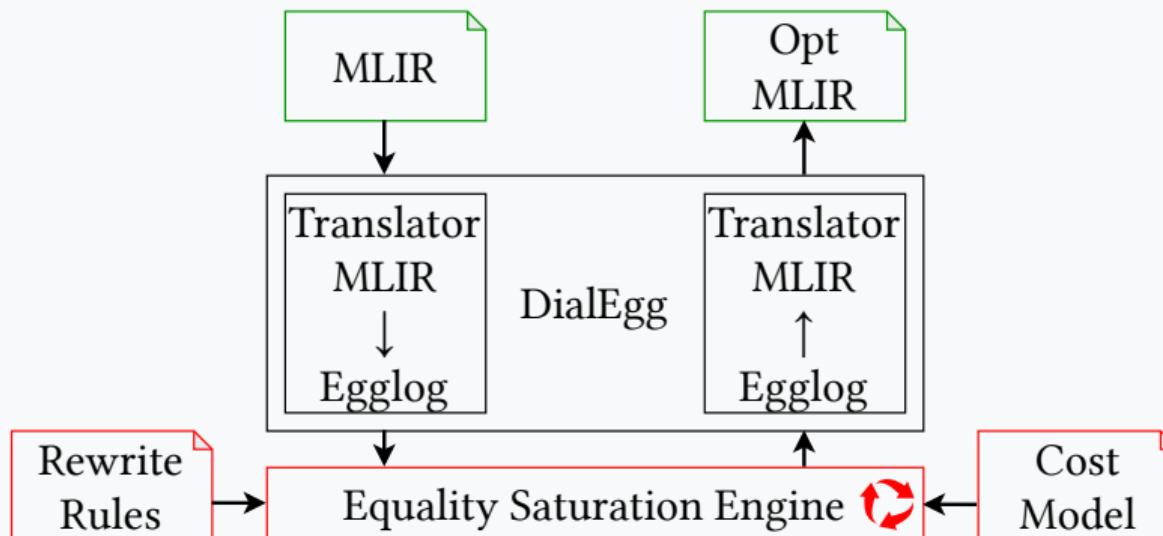
# DialEgg

The first successful attempt to integrate equality saturation with MLIR across abstraction layers.

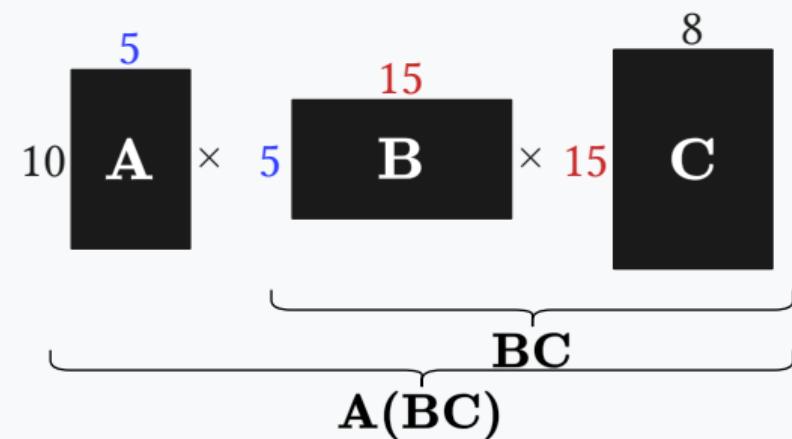
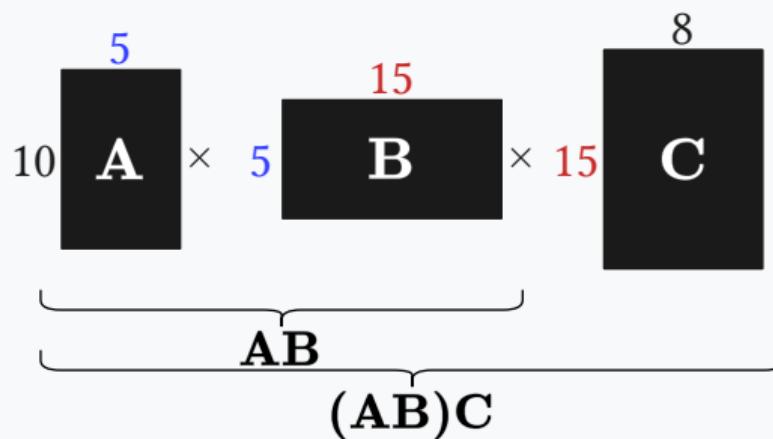


# DialEgg

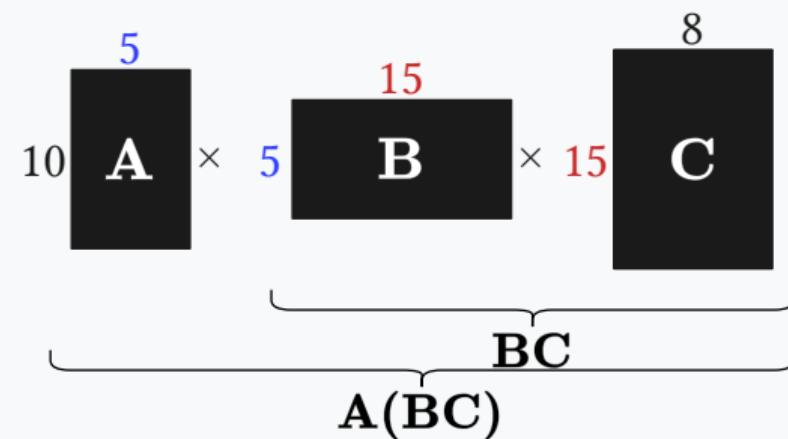
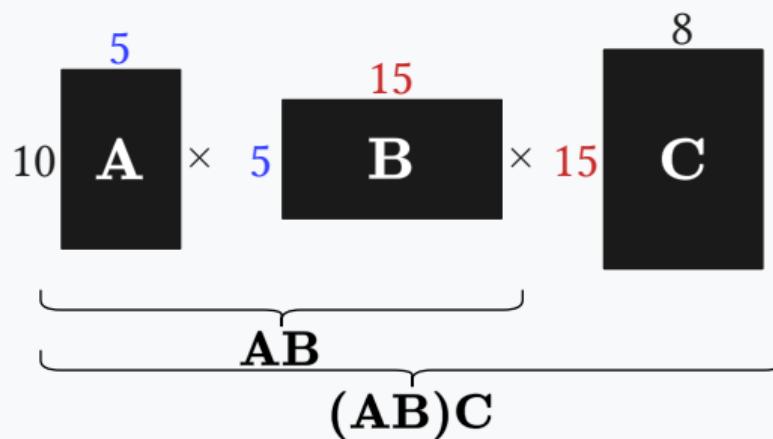
The first successful attempt to integrate equality saturation with MLIR across abstraction layers.



# MatMul is Associative: $(AB)C = A(BC)$

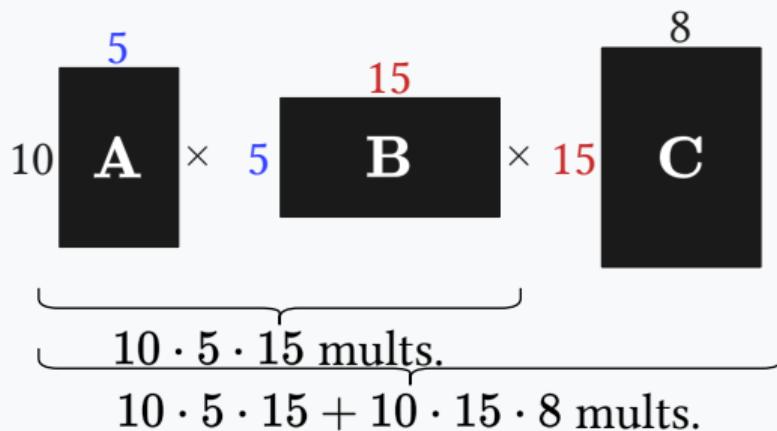


# MatMul is Associative: $(AB)C = A(BC)$

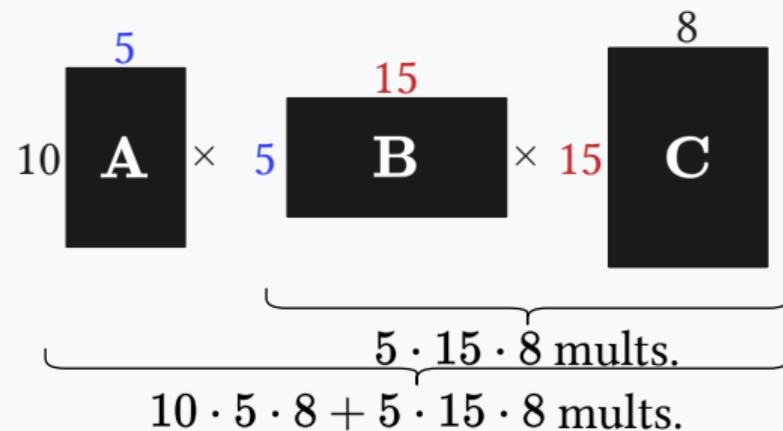


How do we find out which one is more efficient?

# The Order of Operations Matters



$$\text{cost } \mathbf{A}(\mathbf{B}\mathbf{C}) = 1950 \text{ mults.}$$



$$\text{cost } (\mathbf{A}\mathbf{B})\mathbf{C} = 1000 \text{ mults.}$$

# Associativity Rewrite

How we want to express the rewrite  $(XY)Z \Leftrightarrow X(YZ)$ :

```
1 cost (MatMul X Y) = (nrows X) * (ncols X) * (ncols Y)
2 rewrite (MatMul (MatMul X Y) Z) <=> (MatMul X (MatMul Y Z))
```

Repeatedly apply the rewrite to the AST until a fixed-point is reached.

# Associativity Rewrite

How we want to express the rewrite  $(XY)Z \Leftrightarrow X(YZ)$ :

```
1 cost (MatMul X Y) = (nrows X) * (ncols X) * (ncols Y)
2 rewrite (MatMul (MatMul X Y) Z) <=> (MatMul X (MatMul Y Z))
```

Repeatedly apply the rewrite to the AST until a fixed-point is reached.

# Associativity Rewrite

How we want to express the rewrite  $(XY)Z \Leftrightarrow X(YZ)$ :

```
1 cost (MatMul X Y) = (nrows X) * (ncols X) * (ncols Y)
2 rewrite (MatMul (MatMul X Y) Z) <=> (MatMul X (MatMul Y Z))
```

Repeatedly apply the rewrite to the AST until a fixed-point is reached.

# Associativity Rewrite in MLIR

```
class MMAssociateRewritePattern : public OpRewritePattern<MatmulOp> {
public:
    MMAssociateRewritePattern(MLIRContext* context) : OpRewritePattern<MatmulOp>(context) {}

    LogicalResult matchAndRewrite(MatmulOp op, PatternRewriter& rewriter) const override {
        Location loc = op.getLoc();
        Value lhs = op.getOperand(0);
        Value rhs = op.getOperand(1);
        Operation* lhsOp = lhs.getDefiningOp();
        Operation* rhsOp = rhs.getDefiningOp();

        Value x, y, z;
        bool is_x_yz = false, is_xy_z = false;
        if (lhsOp != nullptr && isa<MatmulOp>(lhsOp)) {
            is_xy_z = true;
            x = lhsOp->getOperand(0); // a*b
            y = lhsOp->getOperand(1); // b*c
            z = rhs;                 // c*d
        } else if (rhsOp != nullptr && isa<MatmulOp>(rhsOp)) {
            is_x_yz = true;
            x = lhs;                  // a*b
            y = rhsOp->getOperand(0); // b*c
            z = rhsOp->getOperand(1); // c*d
        } else {
            return failure(); // this is just a normal matmul (XY)
        }
    }
}
```

Pattern matcher

# Associativity Rewrite in MLIR

```
class MMAssociateRewritePattern : public OpRewritePattern<MatmulOp> {
public:
    MMAssociateRewritePattern(MLIRContext* context) : OpRewritePattern<MatmulOp>(context) {}

    RankedTensorType xType = cast<RankedTensorType>(x.getType());
    RankedTensorType yType = cast<RankedTensorType>(y.getType());
    RankedTensorType zType = cast<RankedTensorType>(z.getType());
    Type opType = xType.getElementType();

    if (xType.getNumDynamicDims() > 0 || yType.getNumDynamicDims() > 0 || zType.getNumDynamicDims() > 0) {
        return failure();
    }

    int64_t a = xType.getShape()[0];
    int64_t b = xType.getShape()[1];
    int64_t c = yType.getShape()[1];
    int64_t d = zType.getShape()[1];

    int64_t x_yzCost = b * c * d + a * b * d;
    int64_t xy_zCost = a * b * c + a * c * d;

    } else {
        return failure(); // this is just a normal matmul (XY)
    }
}
```

Cost function

Pattern matcher

# Associativity Rewrite in MLIR

```
class MMAssociateRewritePattern : public OpRewritePattern<MatmulOp> {
public:
    if (xy_zCost < x_yzCost && is_x_yz) { // (XY)Z is better than X(YZ)
        Type xyType = RankedTensorType::get({a, c}, opType);
        Value xyInit = rewriter.create<EmptyOp>(loc, xyType, ValueRange {});
        MatmulOp xy = rewriter.create<MatmulOp>(loc, xyType, ValueRange {x, y}, xyInit);
        Type xyzType = RankedTensorType::get({a, d}, opType);
        Value xyzInit = rewriter.create<EmptyOp>(loc, xyzType, ValueRange {});
        MatmulOp xyz = rewriter.create<MatmulOp>(loc, xyzType, ValueRange {xy.getResult(0), z}, xyzInit);
        rewriter.replaceOp(op, xyz.getResult(0));
    } else if (x_yzCost < xy_zCost && is_xy_z) { // X(YZ) is better than (XY)Z
        Type yzType = RankedTensorType::get({b, d}, opType);
        Value yzInit = rewriter.create<EmptyOp>(loc, yzType, ValueRange {});
        MatmulOp yz = rewriter.create<MatmulOp>(loc, yzType, ValueRange {y, z}, yzInit);
        Type xyzType = RankedTensorType::get({a, d}, opType);
        Value xyzInit = rewriter.create<EmptyOp>(loc, xyzType, ValueRange {});
        MatmulOp xyz = rewriter.create<MatmulOp>(loc, xyzType, ValueRange {x, yz.getResult(0)}, xyzInit);
        rewriter.replaceOp(op, xyz.getResult(0));
    } // if equal, just keep the original

    return success();
};

};
```

Rewriter

# Associativity Rewrite in MLIR

```
class MMAssociateRewritePattern : public OpRewritePattern<MatmulOp> {
public:
    if (xy_zCost < x_yzCost && is_x_yz) { // (XY)Z is better than X(YZ)

class MMAssociatePass : public PassWrapper<MMAssociatePass, OperationPass<FuncOp>> {
public:
   StringRef getArgument() const override { return "mm-assoc"; }
   StringRef getDescription() const override { return "Associates matmul operations." }

    void runOnOperation() override {
        FuncOp func = getOperation();
        MLIRContext* context = &getContext();
        RewritePatternSet patterns(context);
        patterns.add<MMAssociateRewritePattern>(context);

        GreedyRewriteConfig config = GreedyRewriteConfig()
            .setStrictness(GreedyRewriteStrictness::ExistingAndNewOps)
            .setUseTopDownTraversal();
        if (failed(applyPatternsAndFoldGreedily(func, std::move(patterns), config))) {
            signalPassFailure();
        }
    }
};
```

Pass

# Associativity Rewrite in Egglog with DialEgg

```
(rule [(= A (Value ?id ?t))] [(set (type-of A) ?t)])  
(rule [(= A (linalg_matmul ?x ?y ?o ?t))] [(set (type-of A) ?t)])  
  
(function nrows (Type) i64)  
(function ncols (Type) i64)  
(rule [(= ?t (RankedTensor ?shape ?tp))]  
  [(set (nrows ?t) (vec-get ?shape 0))  
   | (set (ncols ?t) (vec-get ?shape 1))])  
  
(rule ; cost[XY] = nrows(X) * ncols(X) * ncols(Y)  
  [(= xy linalg_matmul ?x ?y ?t)]  
  [(cost xy (* (* (nrows (type-of ?x)) (ncols (type-of ?x)))  
              (ncols (type-of ?y))))])  
  
(rule ; (xy)z -> x(yz)  
  [(= xy (linalg_matmul ?x ?y (RankedTensor ?_ ?t)))  
   (= xy_z (linalg_matmul xy ?z ?xyz_t))]  
  
  [(let yz (linalg_matmul ?y ?z (RankedTensor (vec-of (nrows (type-of ?y)) (ncols (type-of ?z))) ?t)))  
   (let x_yz (linalg_matmul ?x yz ?xyz_t))  
   (union xy_z x_yz)])]
```

# Associativity Rewrite in Egglog with DialEgg

```
(rule [(= A (Value ?id ?t))] [(set (type-of A) ?t)])
(rule [(= A (linalg_matmul ?x ?y ?o ?t))] [(set (type-of A) ?t)]))

(function nrows (Type) i64)
(function ncols (Type) i64)
(rule [(= ?t (RankedTensor ?shape ?tp))]
  [(set (nrows ?t) (vec-get ?shape 0))
   | (set (ncols ?t) (vec-get ?shape 1))])

(rule ; cost[XY] = nrows(X) * ncols(X) * ncols(Y)
  [(= xy linalg_matmul ?x ?y ?t)]
  [(cost xy (* (* (nrows (type-of ?x)) (ncols (type-of ?x))) (ncols (type-of ?y))))])

(rule ; (xy)z -> x(yz)
  [(= xy (linalg_matmul ?x ?y (RankedTensor ?_ ?t)))
   (= xy_z (linalg_matmul xy ?z ?xyz_t))]

  [(let yz (linalg_matmul ?y ?z (RankedTensor (vec-of (nrows (type-of ?y)) (ncols (type-of ?z))) ?t)))
   (let x_yz (linalg_matmul ?x yz ?xyz_t))
   (union xy_z x_yz)])]
```

Cost function

# Associativity Rewrite in Egglog with DialEgg

```
(rule [(= A (Value ?id ?t))] [(set (type-of A) ?t)])
(rule [(= A (linalg_matmul ?x ?y ?o ?t))] [(set (type-of A) ?t)])
```

Cost function

```
(function nrows (Type) i64)
(function ncols (Type) i64)
(rule [(= ?t (RankedTensor ?shape ?tp))]
  [(set (nrows ?t) (vec-get ?shape 0))
   | (set (ncols ?t) (vec-get ?shape 1))])
```

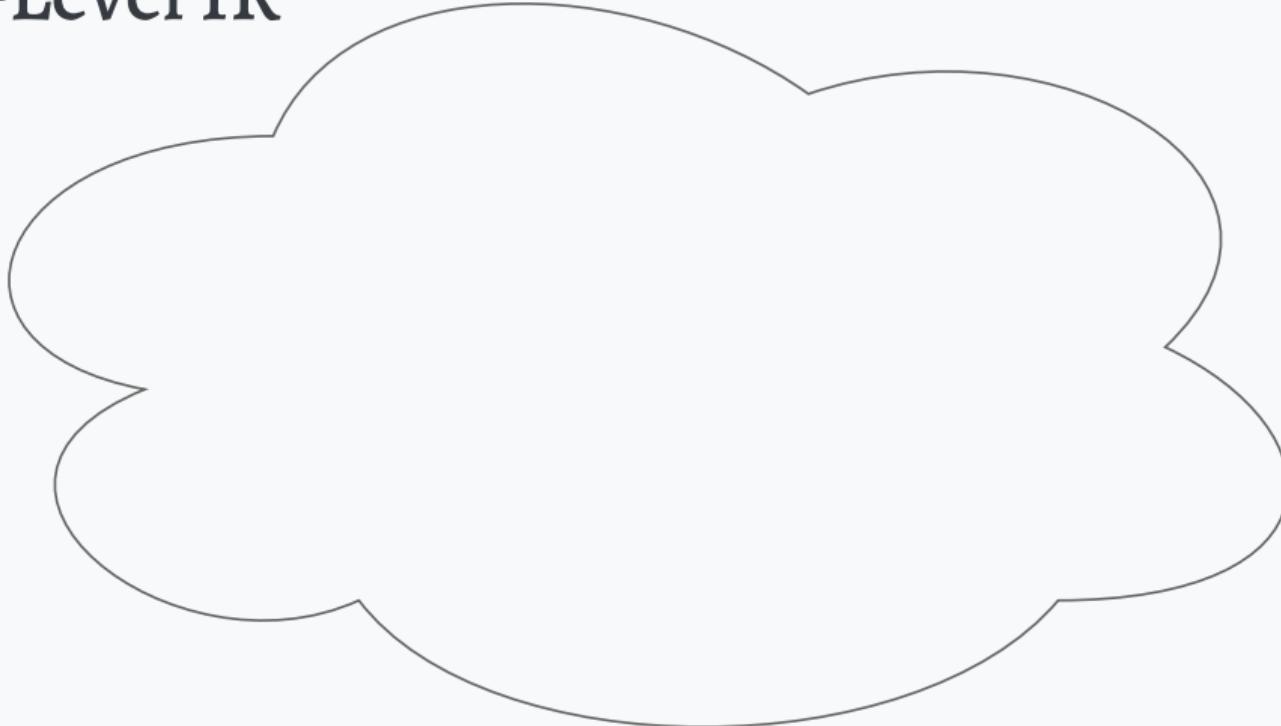
```
(rule ; cost[XY] = nrows(X) * ncols(X) * ncols(Y)
  [= xy linalg_matmul ?x ?y ?t]
  [(cost xy (* (* (nrows (type-of ?x)) (ncols (type-of ?x))) (ncols (type-of ?y))))])
```

```
(rule ; (xy)z -> x(yz)
  [= xy (linalg_matmul ?x ?y (RankedTensor ?_ ?t))]
  [= xy_z (linalg_matmul xy ?z ?xyz_t)])
```

Rewrite rule

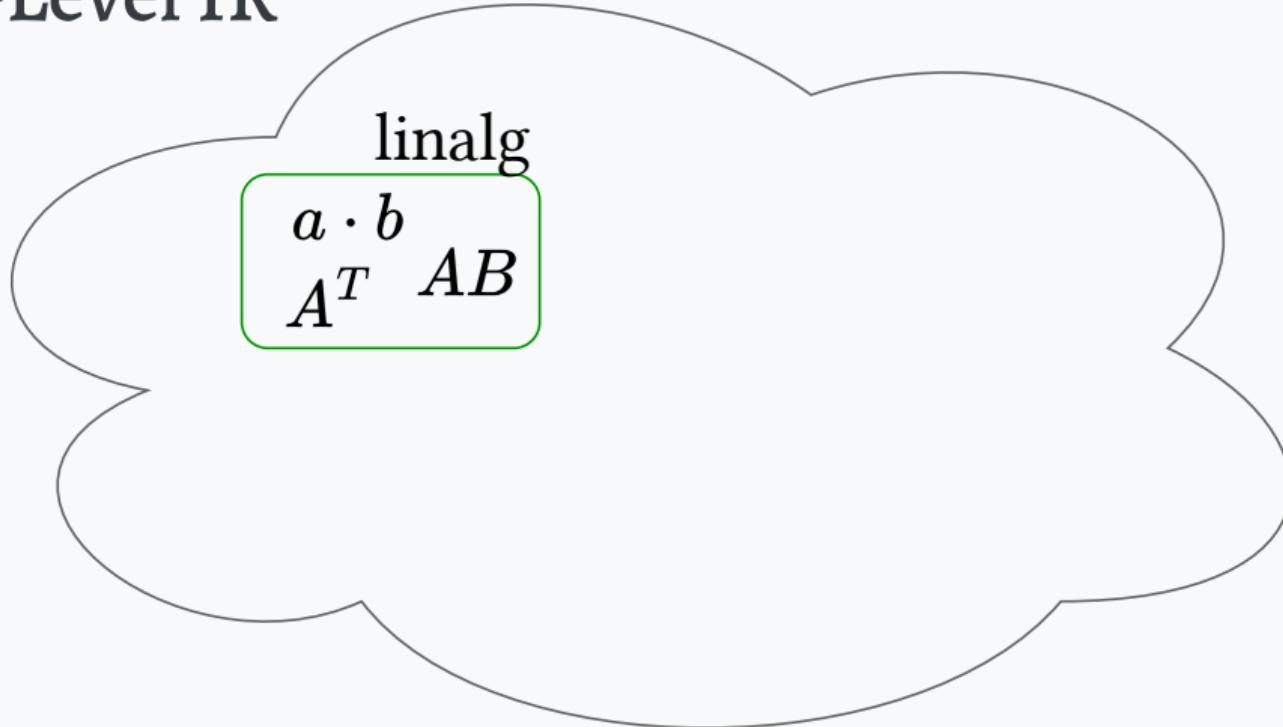
```
[(let yz (linalg_matmul ?y ?z (RankedTensor (vec-of (nrows (type-of ?y)) (ncols (type-of ?z))) ?t)))
  (let x_yz (linalg_matmul ?x yz ?xyz_t))
  (union xy_z x_yz)])
```

# Multi-Level IR



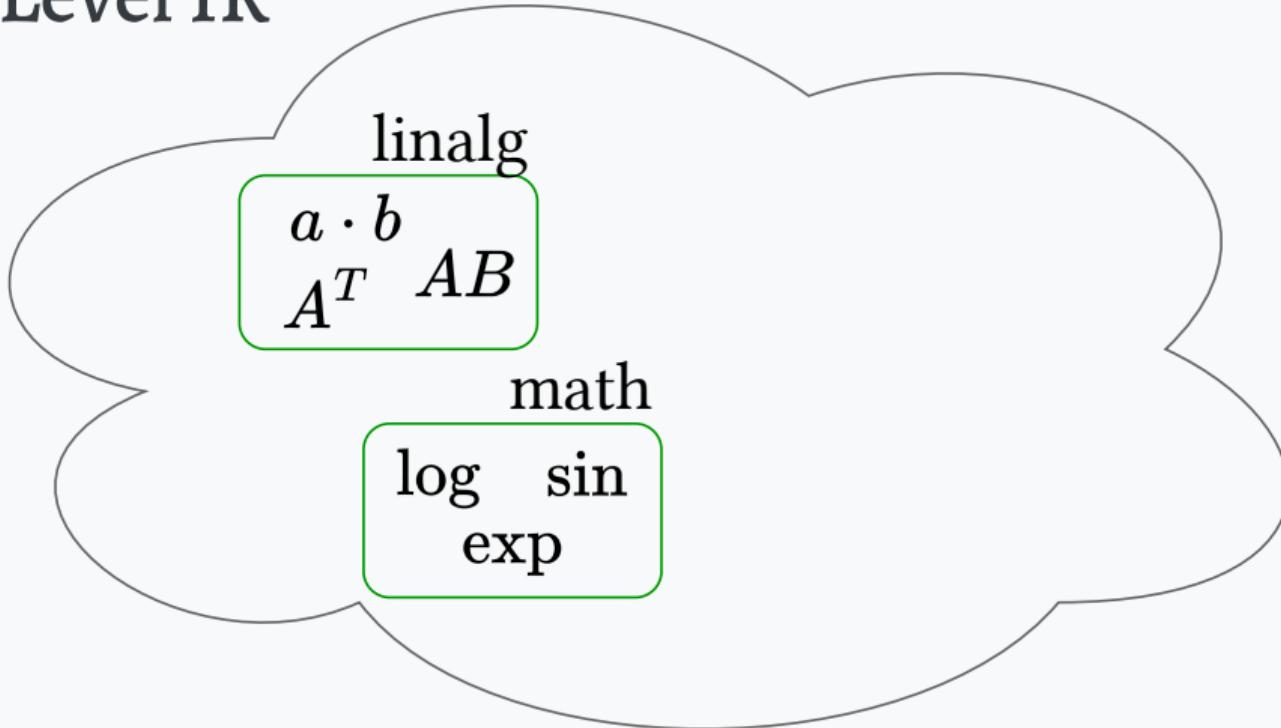
How do we represent many domains and abstraction levels in one IR?

# Multi-Level IR



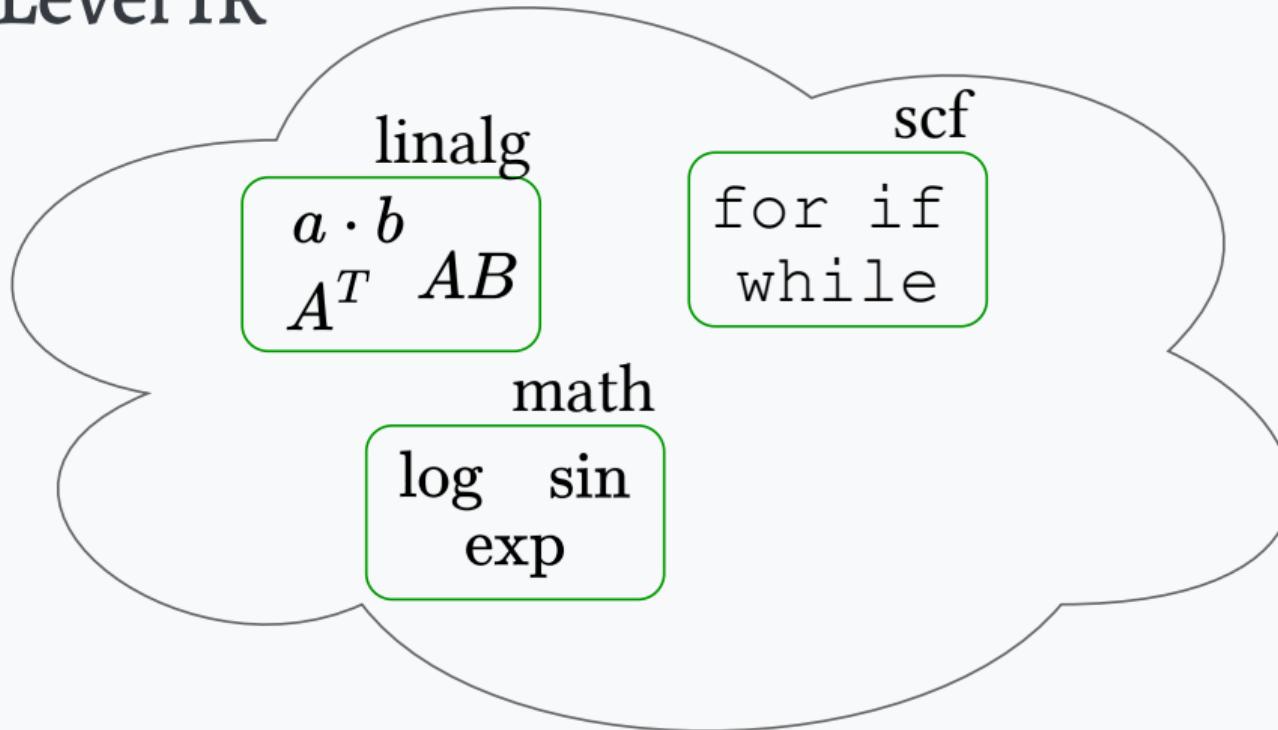
How do we represent many domains and abstraction levels in one IR?

# Multi-Level IR



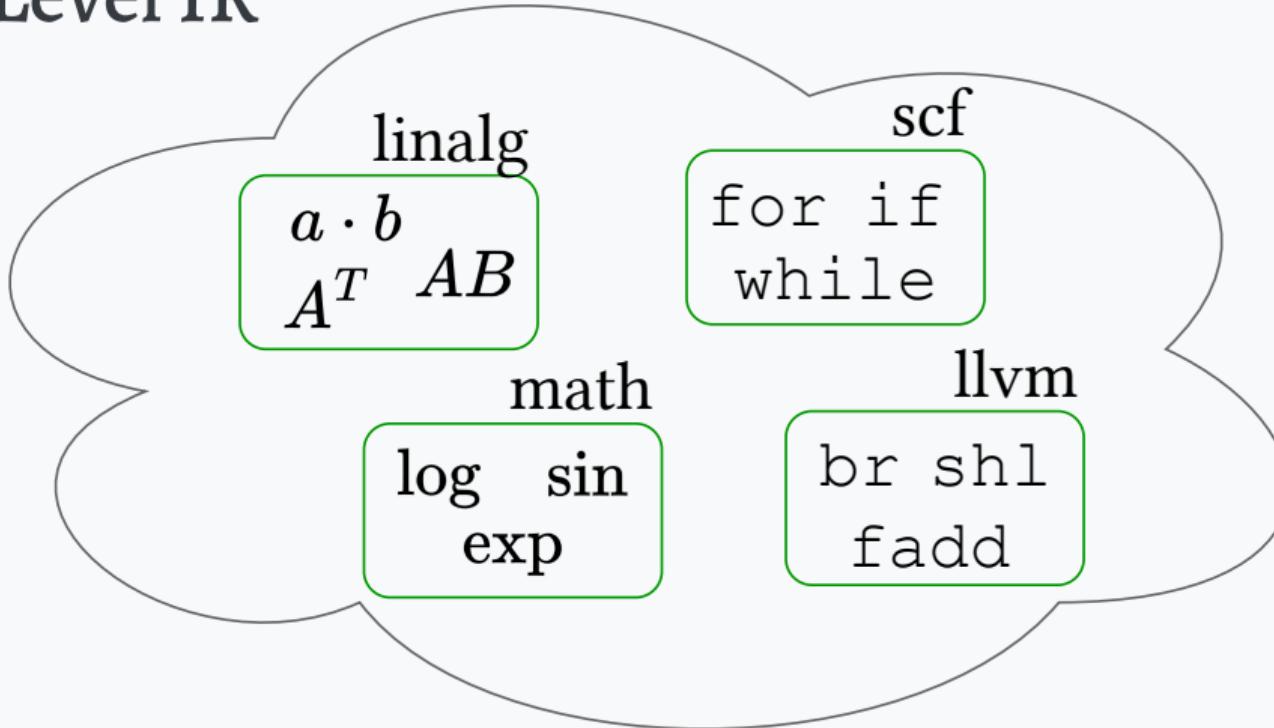
How do we represent many domains and abstraction levels in one IR?

# Multi-Level IR



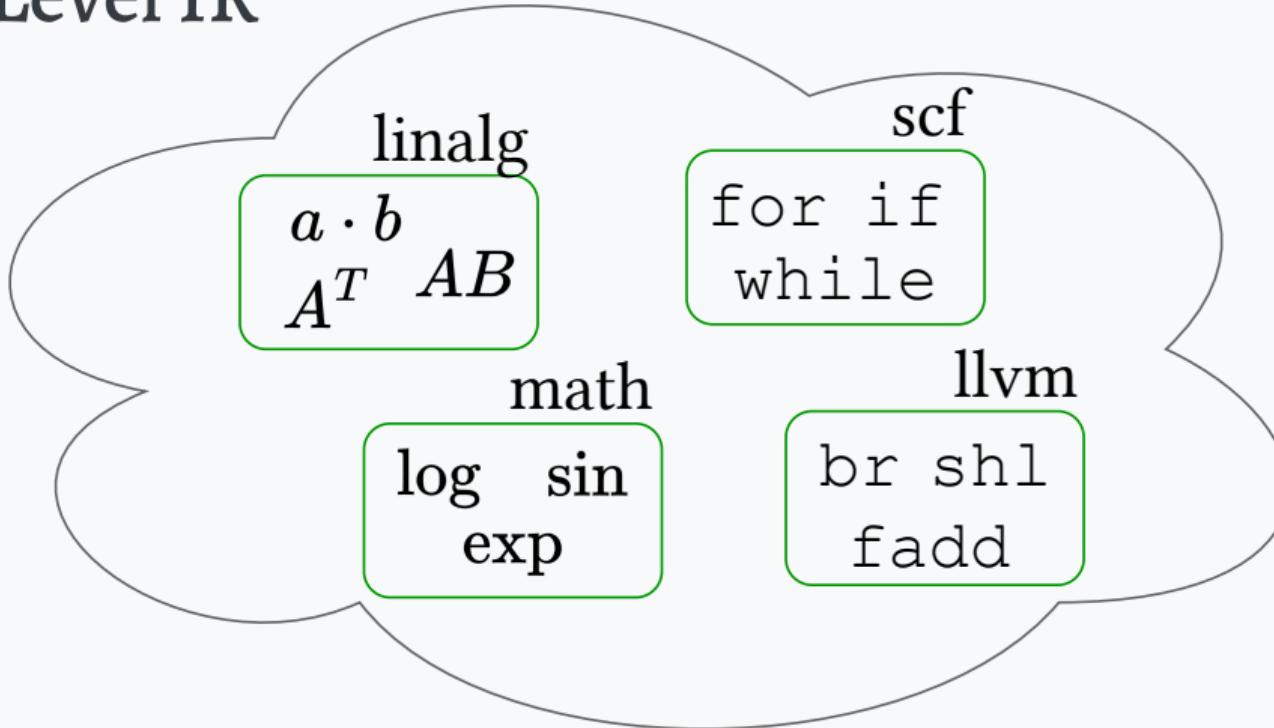
How do we represent many domains and abstraction levels in one IR?

# Multi-Level IR



Abstraction layer  $\approx$  IR level  $\approx$  MLIR dialect.

# Multi-Level IR



And many more: func, arith, gpu, etc.

# MLIR Operations

- Secret sauce:  
recursive operation  
structure

# MLIR Operations

- Secret sauce:  
recursive operation  
structure
  - Operands and results

# MLIR Operations

- Secret sauce:  
recursive operation  
structure
  - Operands and results
  - Attributes

# MLIR Operations

- Secret sauce:  
recursive operation  
structure
  - Operands and results
  - Attributes
  - Regions
    - Blocks (more ops!)

# MLIR Operations

Function to print a vector of size 5.

- Secret sauce:  
recursive operation  
structure
  - Operands and results
  - Attributes
  - Regions
    - Blocks (more ops!)

# MLIR Operations

- Secret sauce:  
recursive operation  
structure
  - Operands and results
  - Attributes
  - Regions
    - Blocks (more ops!)

Function to print a vector of size 5.

```
func.func {sym="tensor_print",
           type=tensor<5xi64> -> () }  
{  
}  
}
```

# MLIR Operations

- Secret sauce:  
recursive operation  
structure
  - Operands and results
  - Attributes
  - Regions
    - Blocks (more ops!)

Function to print a vector of size 5.

```
func.func {sym="tensor_print",
           type=tensor<5xi64> -> () }
{
    %c0 = arith.constant {value=0}: index
    %c5 = arith.constant {value=5}: index
}
```

# MLIR Operations

- Secret sauce:  
recursive operation  
structure
  - Operands and results
  - Attributes
  - Regions
    - Blocks (more ops!)

Function to print a vector of size 5.

```
func.func {sym="tensor_print",
           type=tensor<5xi64> -> () }
{
    %c0 = arith.constant {value=0}: index
    %c5 = arith.constant {value=5}: index
    scf.for %i = %c0 to %c5
    {
        %c1 = arith.constant {value=1}: index
        %c2 = arith.constant {value=2}: index
        %c3 = arith.constant {value=3}: index
        %c4 = arith.constant {value=4}: index
    }
}
```

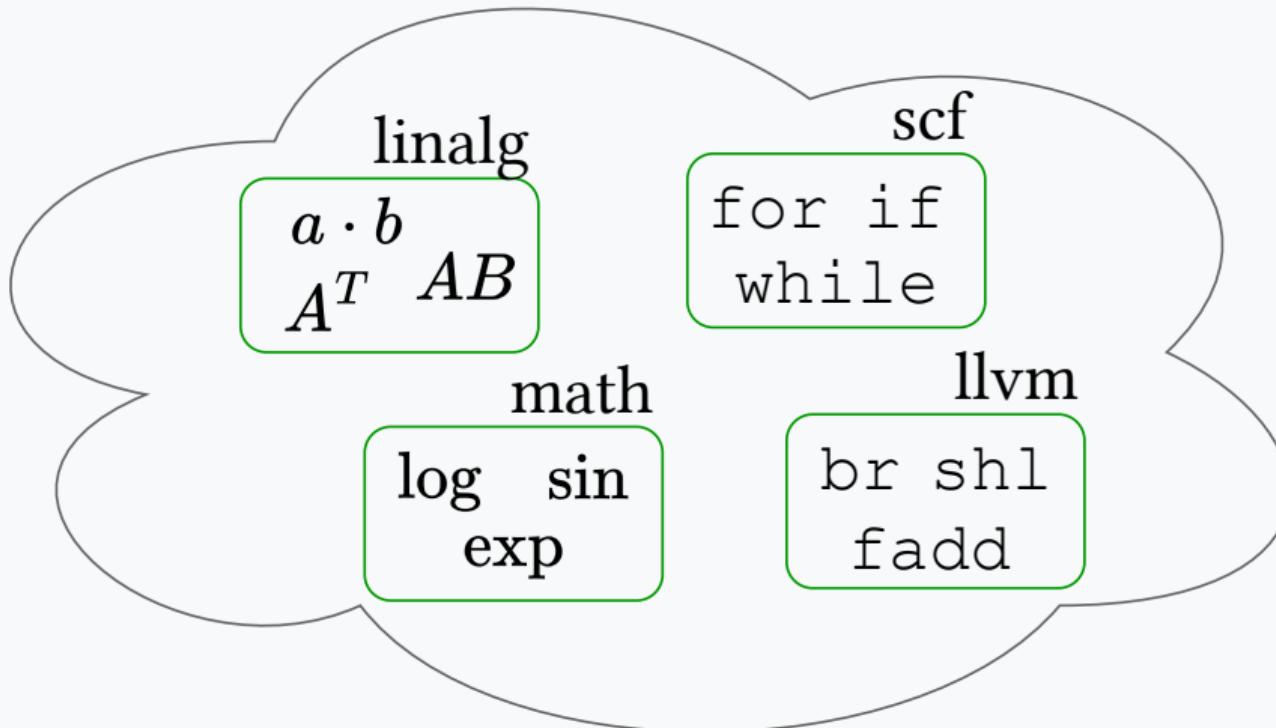
# MLIR Operations

- Secret sauce:  
recursive operation  
structure
  - Operands and results
  - Attributes
  - Regions
    - Blocks (more ops!)

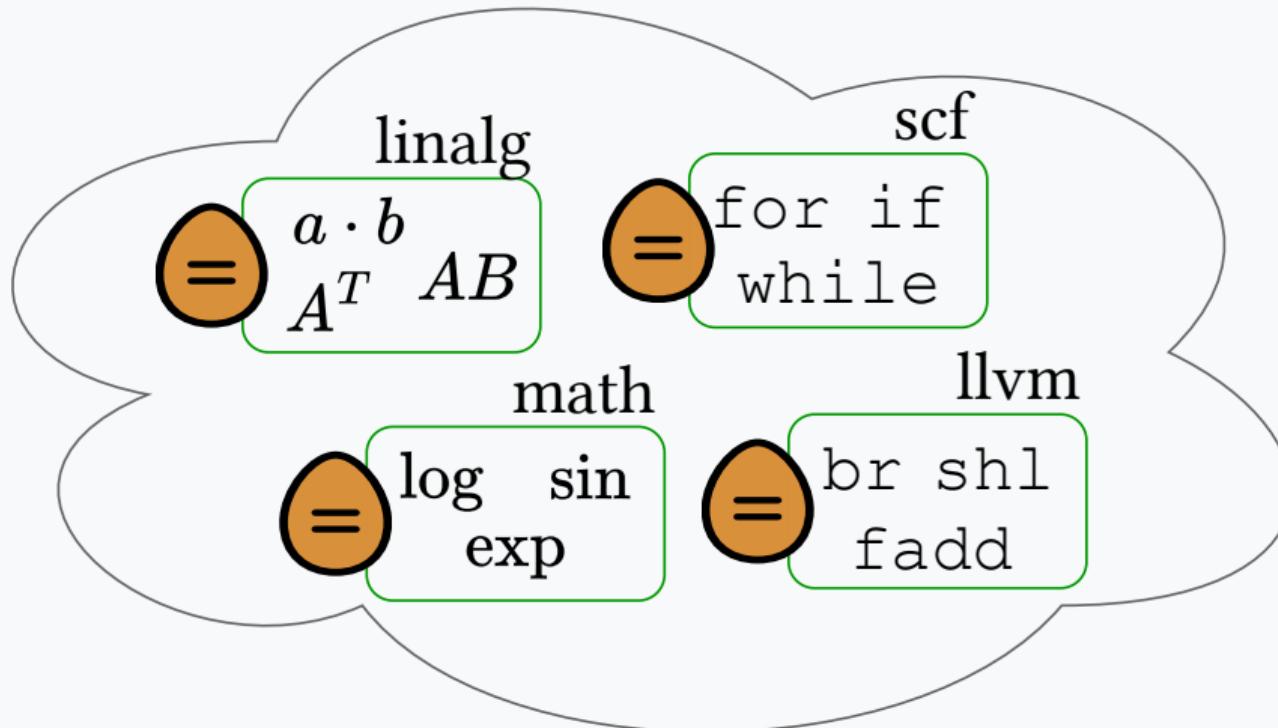
Function to print a vector of size 5.

```
func.func {sym="tensor_print",
           type=tensor<5xi64> -> () }
{
    %c0 = arith.constant {value=0}: index
    %c5 = arith.constant {value=5}: index
    scf.for %i = %c0 to %c5
    {
        %e = tensor.extract %t[%i]: tensor<5xi64>
        func.call(%e) {callee="print_i64"}: (i64)->()
    }
}
```

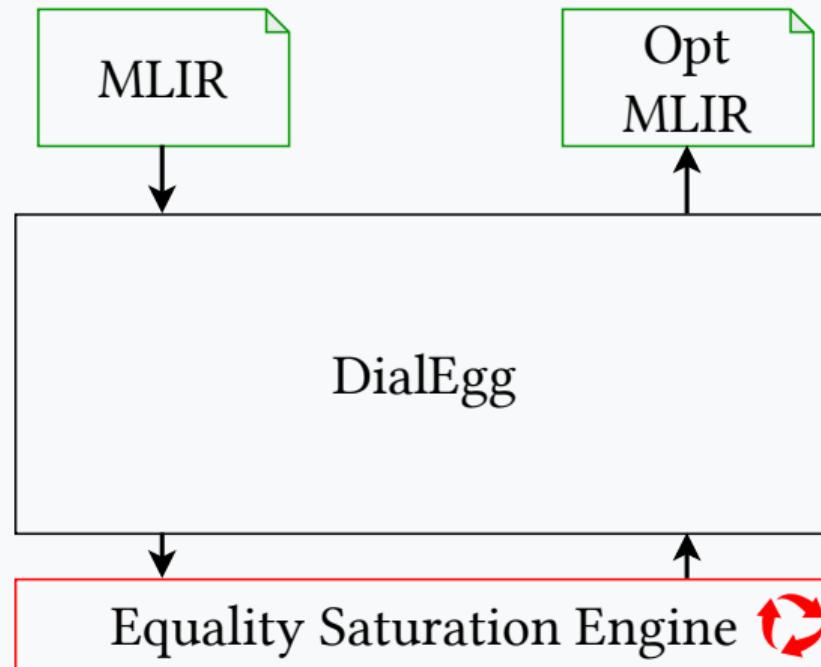
# DialEgg: A Dialect-Agnostic MLIR Optimizer



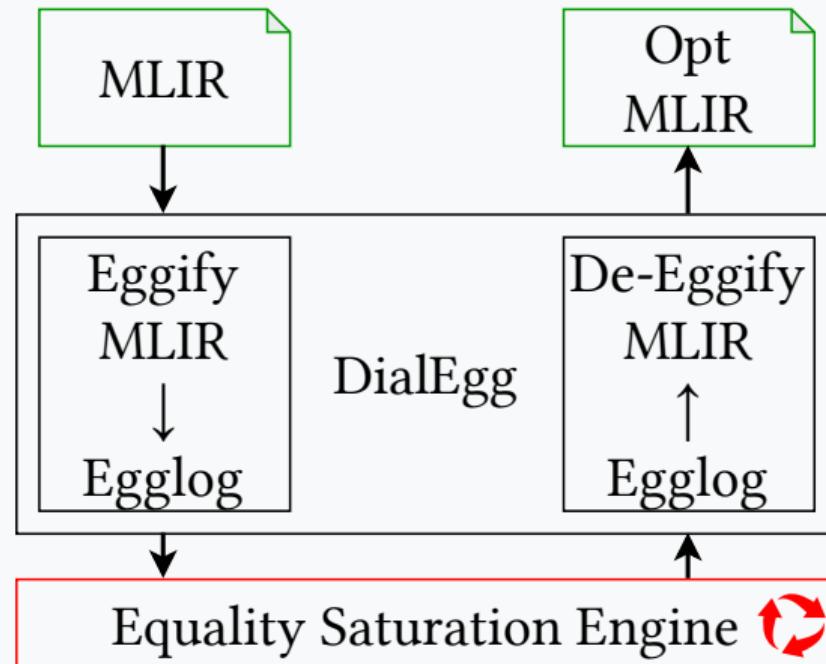
# DialEgg: A Dialect-Agnostic MLIR Optimizer



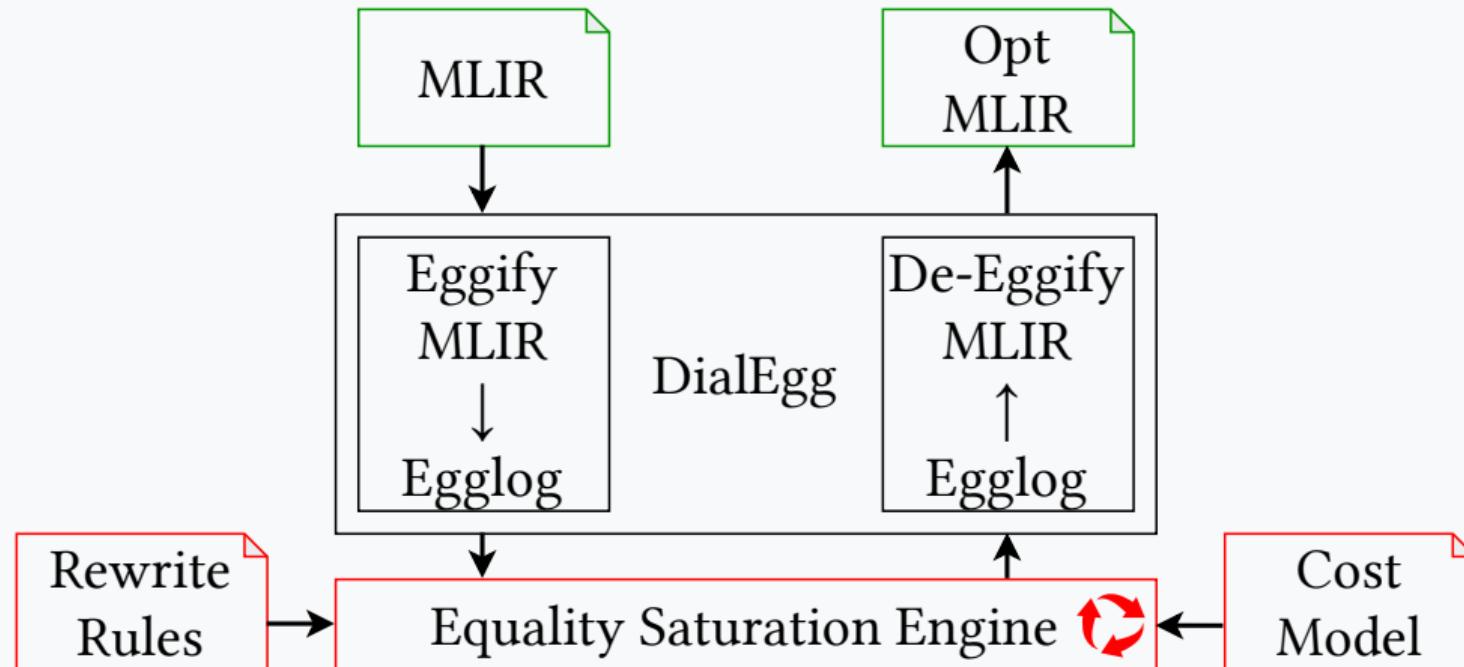
# DialEgg



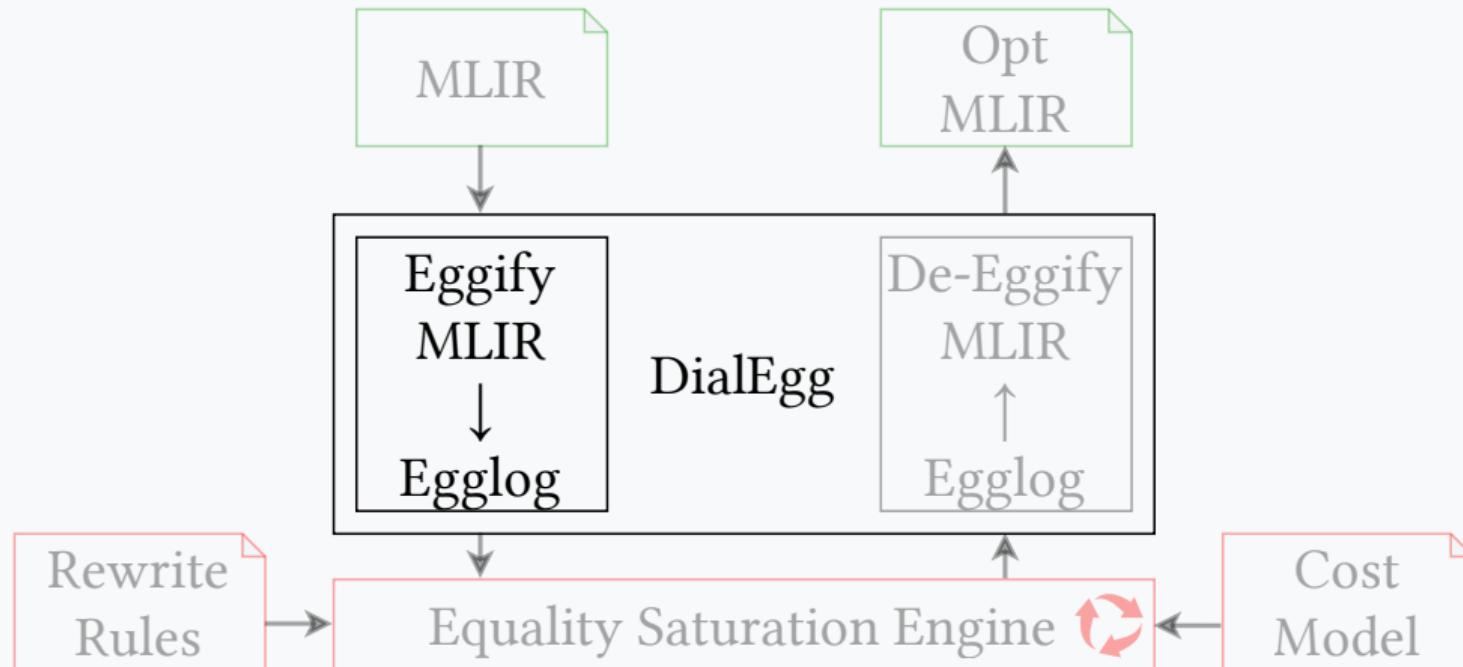
# DialEgg



# DialEgg



# MLIR → Egglog



# MLIR → Egglog: Types

MLIR types are eggified to variants of the Type ADT in Egglog.

```
1 (datatype Type
2   (F32)
3   (I64)
4   (Tensor IntVec Type)
5   (OpaqueType String String) ; (type, value)
6   ...
7 )
```

# MLIR → Egglog: Types

MLIR types are eggified to variants of the Type ADT in Egglog.

```
1 (datatype Type
2   (F32)
3   (I64)
4   (Tensor IntVec Type)
5   (OpaqueType String String) ; (type, value)
6   ...
7 )
```

Tensor Type in Egglog

```
1 (Tensor (vec-of 2 3) (I64)) ; tensor<2x3xi64>
```

# MLIR → Egglog: Attributes

MLIR attributes are eggified to variants of the Attr ADT in Egglog.

```
1 (datatype Attr
2   (IAttr i64 Type)
3   (DenseIntElementsAttr IntVec Type)
4   (OpaqueAttr String String) ; (name, value)
5   ...
6 )
7 (datatype NamedAttrT (NamedAttr String Attr))
```

# MLIR → Egglog: Attributes

MLIR attributes are eggified to variants of the Attr ADT in Egglog.

```
1 (datatype Attr
2   (IAttr i64 Type)
3   (DenseIntElementsAttr IntVec Type)
4   (OpaqueAttr String String) ; (name, value)
5   ...
6 )
7 (datatype NamedAttrT (NamedAttr String Attr))
```

Constant operation in MLIR

```
1 arith.const {value=5: i64} : i64
```

```
1 (NamedAttr "value" (IAttr 5 (I64)))
```

# MLIR → Egglog: Operations

Operations are eggified piece by piece to variants of the Op ADT.

```
1 (datatype Op
2   (Value i64 Type) ; (id, type)
3   (arith_const NamedAttrT Type)
4   (arith_addi Op Op Type)
5   (linalg_matmul Op Op Type)
6   (scf_if Op Region Region Type)
7   ...
8 )
```

# MLIR → Egglog: Operations

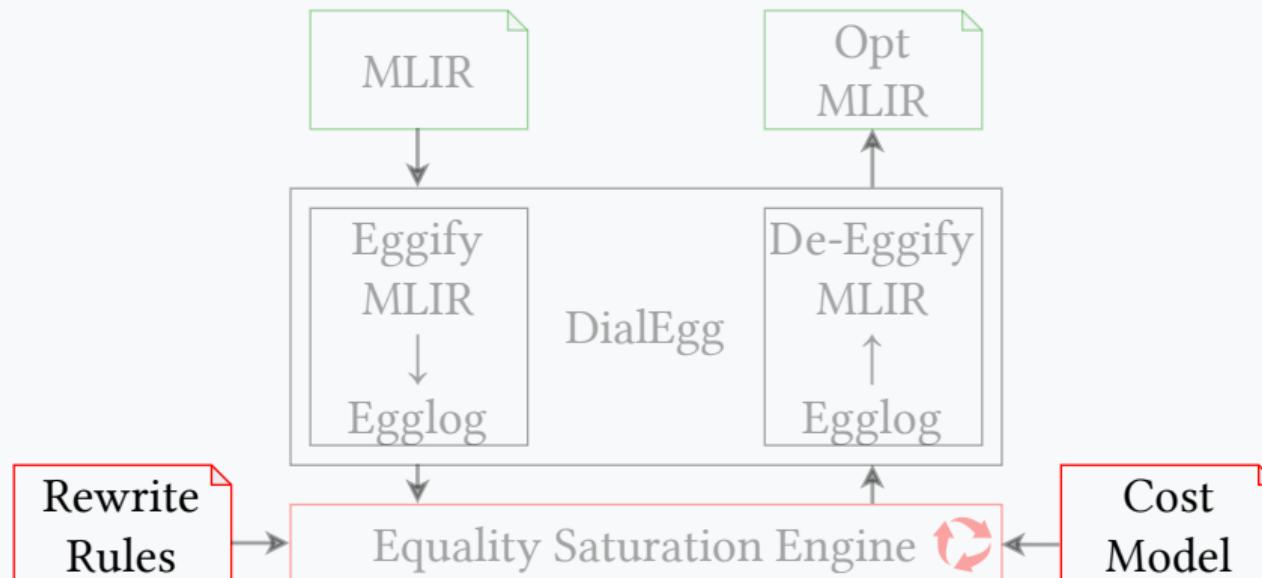
Operations are eggified piece by piece to variants of the Op ADT.

```
1 (datatype Op
2   (Value i64 Type) ; (id, type)
3   (arith_const NamedAttrT Type)
4   (arith_addi Op Op Type)
5   (linalg_matmul Op Op Type)
6   (scf_if Op Region Region Type)
7   ...
8 )
```

MatMul Operation in Egglog

```
1 (linalg_matmul A B (Tensor (vec-of 20 10) (I64)))
```

# MM Associativity Rewrite Rule in Egglog



# MM Associativity Rewrite Rule in Egglog

$$(XY)Z \Rightarrow X(YZ)$$

where  $\mathbf{X}: a \times b$ ,  $\mathbf{Y}: b \times c$ , and  $\mathbf{Z}: c \times d$

```
1 (rule
2   [ (= XY (linalg_matmul ?X ?Y (Tensor ?_ ?t)))
3     (= XY_Z (linalg_matmul XY ?Z ?XYZ_t))
4     (= b (nrows type-of ?Y))
5     (= d (ncols type-of ?Z)) ]
6
7   [(let YZ (linalg_matmul ?Y ?Z (Tensor (vec-of b d) ?t)))
8     (let X_YZ (linalg_matmul ?X YZ ?XYZ_t))
9     (union XY_Z X_YZ))]
10 )
```

# MM Associativity Rewrite Rule in Egglog

$$(XY)Z \Rightarrow X(YZ)$$

where  $\mathbf{X}: a \times b$ ,  $\mathbf{Y}: b \times c$ , and  $\mathbf{Z}: c \times d$

```
1 (rule
2   [= XY (linalg_matmul ?X ?Y (Tensor ?_ ?t)))
3   [= XY_Z (linalg_matmul XY ?Z ?XYZ_t))
4   [= b (nrows type-of ?Y))
5   [= d (ncols type-of ?Z))]
6
7   [(let YZ (linalg_matmul ?Y ?Z (Tensor (vec-of b d) ?t)))
8     (let X_YZ (linalg_matmul ?X YZ ?XYZ_t))
9     (union XY_Z X_YZ))]
10 )
```

# MM Associativity Cost Model in Egglog

$$\text{cost } \mathbf{XY} = abc$$

where  $\mathbf{X}: a \times b$  and  $\mathbf{Y}: b \times c$

```
1 (rule
2   [(= XY (linalg_matmul ?X ?Y ?XY_t))
3    (= a (nrows type-of ?X))
4    (= b (ncols type-of ?X))
5    (= c (ncols type-of ?Y))])
6
7   [((cost XY (* (* a b) c)))]
8 )
```

# MM Associativity Cost Model in Egglog

$$\text{cost } \mathbf{XY} = abc$$

where  $\mathbf{X}: a \times b$  and  $\mathbf{Y}: b \times c$

```
1 (rule
2   [ (= XY (linalg_matmul ?X ?Y ?XY_t))
3     (= a (nrows type-of ?X))
4     (= b (ncols type-of ?X))
5     (= c (ncols type-of ?Y)) ]
6
7   [(cost XY (* (* a b) c)) ]
8 )
```

# Optimized 3MM

```
1 func.func @3mm(%A,%B,%C,%D) {  
2   %AB =linalg.matmul %A    %B  
3   %ABC =linalg.matmul %AB    %C  
4   %ABCD=linalg.matmul %ABC %D  
5   func.return %ABCD  
6 }
```

cost  $((\mathbf{AB}) \mathbf{C}) \mathbf{D} = 10,000$  mults.

# Optimized 3MM

```
1 func.func @3mm(%A,%B,%C,%D) {  
2   %AB =linalg.matmul %A %B  
3   %ABC =linalg.matmul %AB %C  
4   %ABCD=linalg.matmul %ABC %D  
5   func.return %ABCD  
6 }
```

cost  $((AB)C)D = 10,000$  mults.

```
1 func.func @3mm(%A,%B,%C,%D) {  
2   %CD =linalg.matmul %C %D  
3   %BCD =linalg.matmul %B %CD  
4   %ABCD=linalg.matmul %A %BCD  
5   func.return %ABCD  
6 }
```

cost  $A(B(CD)) = 5,000$  mults.

# Case Study: Polynomial Evaluation

Reduce from  $O(n^2)$  to  $O(n)$  multiplications using Horner's method.

$$\begin{aligned}P(x) &= a_0 + a_1x + a_2x^2 + \cdots + a_{n-1}x^{n-1} + a_nx^n \\&= a_0 + x(a_1 + x(a_2 + \cdots + x(a_{n-1} + xa_n)))\end{aligned}$$

# Case Study: Polynomial Evaluation

Reduce from  $O(n^2)$  to  $O(n)$  multiplications using Horner's method.

$$\begin{aligned}P(x) &= a_0 + a_1x + a_2x^2 + \cdots + a_{n-1}x^{n-1} + a_nx^n \\&= a_0 + x(a_1 + x(a_2 + \cdots + x(a_{n-1} + xa_n)))\end{aligned}$$

- Exponentiation:  $x^0 \Rightarrow 1$  and  $x^n \Leftrightarrow x \cdot x^{n-1}$

# Case Study: Polynomial Evaluation

Reduce from  $O(n^2)$  to  $O(n)$  multiplications using Horner's method.

$$\begin{aligned}P(x) &= a_0 + a_1x + a_2x^2 + \cdots + a_{n-1}x^{n-1} + a_nx^n \\&= a_0 + x(a_1 + x(a_2 + \cdots + x(a_{n-1} + xa_n)))\end{aligned}$$

- Exponentiation:  $x^0 \Rightarrow 1$  and  $x^n \Leftrightarrow x \cdot x^{n-1}$
- Commutativity:  $x + y \Leftrightarrow y + x$  and  $x \cdot y \Leftrightarrow y \cdot x$
- Associativity:  $(x + y) + z \Leftrightarrow x + (y + z)$  and  $(x \cdot y) \cdot z \Leftrightarrow x \cdot (y \cdot z)$
- Distributivity:  $x \cdot (y + z) \Leftrightarrow x \cdot y + x \cdot z$
- Identity:  $x \cdot 1 \Rightarrow x$

# Case Study: Polynomial Evaluation

- Exponentiation:  $x^n \Leftrightarrow x \cdot x^{n-1}$

```
1 (rule
2   [= n1 (arith_const (NamedAttr "value" (IAttr ?n)) ?t))
3     (= x1 (math_ipowi ?x n1 ?t)) ; x^n
4     (>= ?n 1])
5
6   [(let n2 (arith_const (NamedAttr "value" (IAttr (- ?n 1))) ?t))
7     (let x2 (math_ipowi ?x n2 ?t)) ; x^{n-1}
8     (let x3 (arith_muli ?x x2 ?t)) ; x * x^{n-1}
9     (union x1 x3)]
10 )
```

# Compile Time and Scalability

Bench.	#Rules	#Ops	MLIR	Egglog	Egglog	Saturation
			↓↑			
Img Conv	1	29	0.4ms	14.6ms	<0.1ms	
Vec Norm	1	44	0.5ms	21.6ms	<0.1ms	
Poly	8	26	0.5ms	18.9ms	2ms	
3MM	5	8	0.3ms	8.7ms	1ms	
10MM	5	22	0.5ms	14.4ms	4ms	
20MM	5	42	1.0ms	41.3ms	23ms	
40MM	5	82	1.8ms	296.2ms	235ms	
80MM	5	162	7.3ms	4939.3ms	3732ms	

# DialEgg



azizzayed.com

- More case studies in the paper:
  - Image conversion: rewrites with computation.
  - Vector normalization: rewrites matching on MLIR attributes.



DialEgg paper

# MM Associativity Helpers

```
1 (rule ((= ?A (Value ?id ?t))) ((set (type-of ?A) ?t)))
2 (rule ((= ?A (linalg_matmul ?x ?y ?t))) ((set (type-of ?A) ?t)))
3
4 (function nrows (Type) i64)
5 (function ncols (Type) i64)
6 (rule ((= ?t (RankedTensor ?dims ?tp)))
7     ((set (nrows ?t) (vec-get ?dims 0))
8     (set (ncols ?t) (vec-get ?dims 1))))
9
10 (rule ((linalg_matmul ?x ?y ?t)
11     (= a (nrows (type-of ?x))))
12     (= b (ncols (type-of ?x))))
13     (= c (ncols (type-of ?y))))
14
15     ((cost (linalg_matmul ?x ?y ?t) (* (* a b) c))))
```

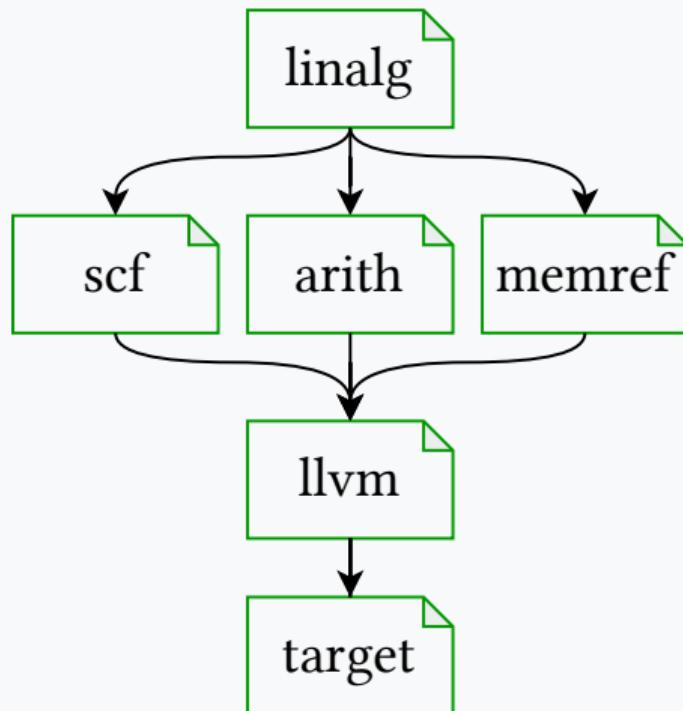
# Poly Rewrite Rules

```
1  (rewrite (arith_addi ?x ?y ?a ?t) (arith_addi ?y ?x ?a ?t))
2  (rewrite (arith_muli ?x ?y ?a ?t) (arith_muli ?y ?x ?a ?t))
3  (rewrite ; (x + y) + z = x + (y + z)
4    (arith_addi (arith_addi ?x ?y ?a ?t) ?z ?a ?t)
5    (arith_addi ?x (arith_addi ?y ?z ?a ?t) ?a ?t))
6  (rewrite ; (x * y) * z = x * (y * z)
7    (arith_muli (arith_muli ?x ?y ?a ?t) ?z ?a ?t)
8    (arith_muli ?x (arith_muli ?y ?z ?a ?t) ?a ?t))
9
10 (rewrite (arith_muli ?x ; x * 1 = x
11   (arith_const (NamedAttr "value" (IAttr 1.0 ?t)) ?t)
12   ?a ?t) ?x)
13 (rewrite (math_ipowi ?x ; x^0 = 1
14   (arith_const (NamedAttr "value" (IAttr 0.0 ?t)) ?t) ?a ?t)
15   (arith_const (NamedAttr "value" (IAttr 1.0 ?t)) ?t))
```

# Poly Rewrite Rules

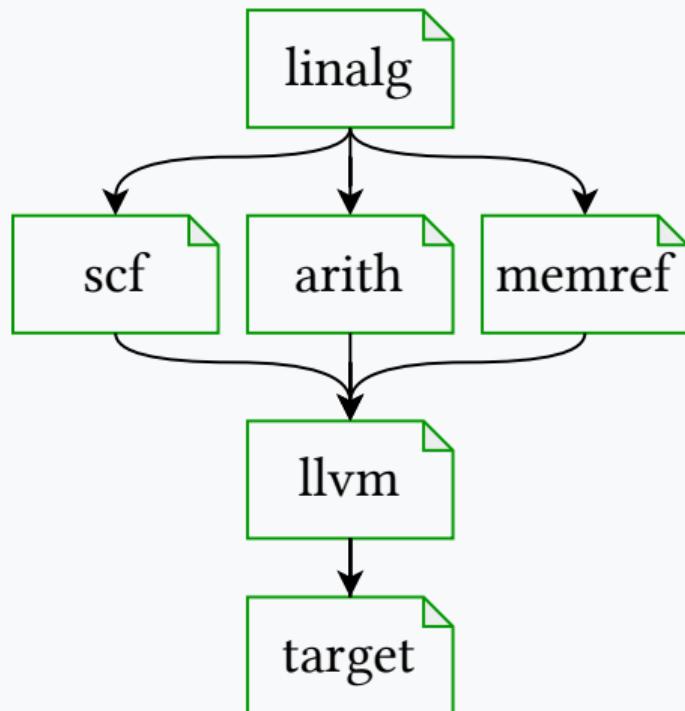
```
1  (rewrite ; mx + nx = x(m + n)
2    (arith_addi (arith_muli ?m ?x ?a ?t) (arith_muli ?n ?x ?a ?t)
3      ?a ?t)
4    (arith_muli ?x (arith_addi ?m ?n ?a ?t) ?a ?t))
5
6  (rule
7   [(= n1 (arith_const (NamedAttr "value" (IAttr ?n)) ?t))
8    (= x1 (math_ipowi ?x n1 ?t)) ; x^n
9    (>= ?n 1)]
10
11  [(let n2 (arith_const (NamedAttr "value" (IAttr (- ?n 1))) ?t))
12   (let x2 (math_ipowi ?x n2 ?t)) ; x^{n-1}
13   (let x3 (arith_muli ?x x2 ?t)) ; x * x^{n-1}
14   (union x1 x3)])
15 )
```

# MLIR Lowering and Code Generation



Lower to a lower level of abstraction.

# MLIR Lowering and Code Generation



Lower to a lower level of abstraction.

```
// X: tensor<5x10xi64>
// Y: tensor<10x3xi64>
%XY = linalg.matmul %X %Y
```

# MLIR Lowering and Code Generation

```
// X: tensor<5x10xi64>
// Y: tensor<10x3xi64>
%XY = linalg.matmul %X %Y
```

# MLIR Lowering and Code Generation

## MatMul with scf.for

```
// X: tensor<5x10xi64>
// Y: tensor<10x3xi64>
%XY = linalg.matmul %X %Y
```

```
%XY = memref.alloc(): memref<5x3xi64>
scf.for %i = %0 to %5 {
    scf.for %j = %0 to %3 {
        scf.for %k = %0 to %10 {
            %X_ik = memref.load %X[%i, %k]
            %Y_kj = memref.load %Y[%k, %j]
            %XY_ij = memref.load %XY[%i, %j]
            %a = arith.muli %X_ik, %Y_kj
            %b = arith.addi %XY_ij, %a
            memref.store %b, %XY[%i, %j]
        }
    }
}
```